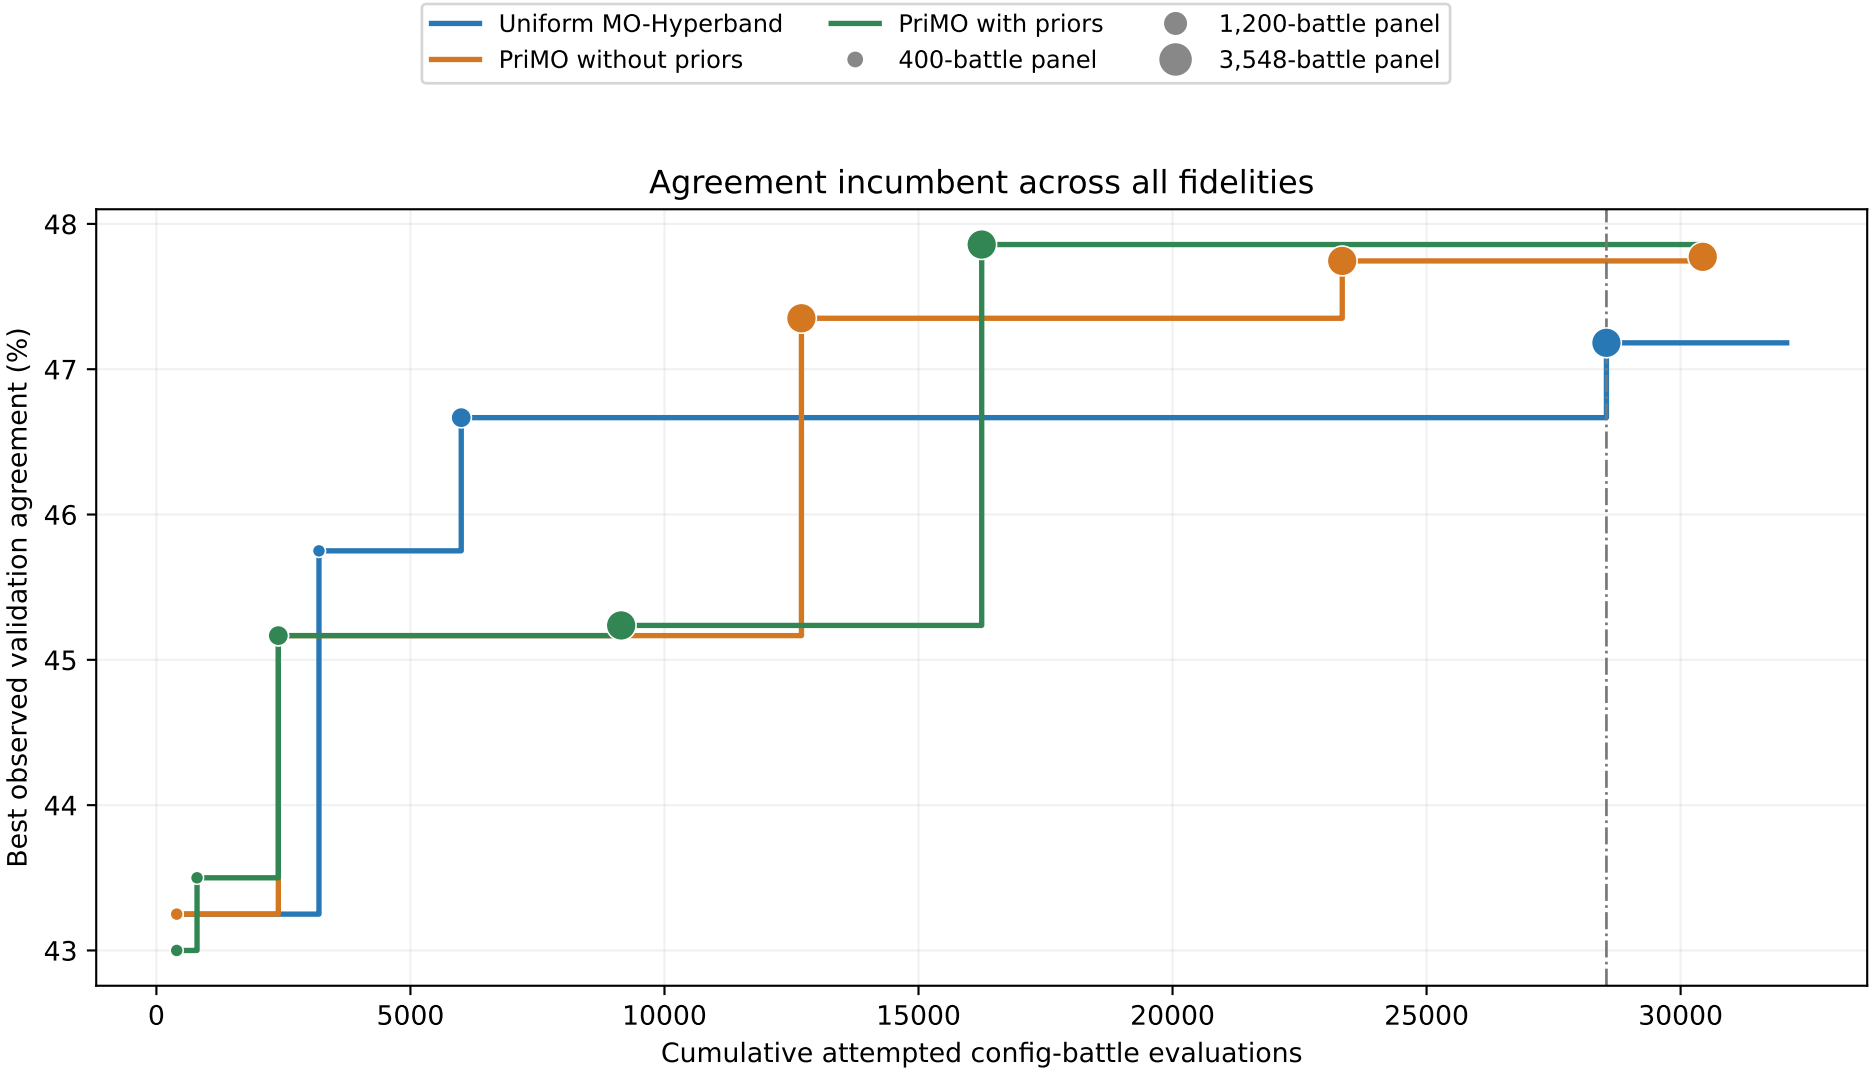

Source: verified trial reports and attempt logs, 14–15 Oct 2026. One best-so-far curve per algorithm, pooling observed scores across all fidelities.
Circle size identifies the panel size of each incumbent-setting evaluation. Dash-dot line: 28,540 attempted config-battles; timeout work is included.
Mixed-fidelity scores use different validation panel sizes; this is an observed-search trajectory, not a full-fidelity performance comparison. One seed per arm.